Test Task — Agentic AI Engineer

Flat Rock Technology

Overview

Build a working system where an email with a file attachment gets picked up by an AI agent and processed — end to end.

Pick a business scenario that makes sense to you. Contracts, reports, purchase orders, onboarding documents — your call. The one thing we ask: don't do invoices. Every automation platform has a ready-made template for that, and we want to see your thinking, not someone else's.

What to build

Email intake — Set up a webhook or trigger that receives an email with a file attachment (PDF, image, CSV — whatever fits your scenario).

Extraction — The agent reads the attached file and pulls out structured data as JSON. You define the schema based on your chosen scenario. The extraction should be accurate and handle real-world messiness — not just a clean, perfectly formatted sample.

RAG — Load a small set of relevant documents into a vector store. The agent should query this knowledge base to enrich or validate what it extracted. For example: cross-referencing an extracted company name against known clients, or checking terms against internal policies. The RAG piece should add real value, not just be tacked on.

Routing & action — Based on what the agent finds, it should do something useful. Clean, high-confidence results get stored somewhere (spreadsheet, database, notification — up to you). Anything the agent isn't sure about gets flagged for human review. Show us you've thought about what happens when things aren't perfect.

Stack

Use whatever you're comfortable with. Any LLM provider (OpenRouter, OpenAI, Anthropic, etc.), any workflow or automation tool (n8n, Make, Activepieces, or your own), any vector store. We care about the result, not the brand names.

Deliverables

- A public GitHub repo with your code and any workflow exports.
- A short screen recording (Loom, OBS, whatever) walking through 3 different file attachments being processed. We want to see it run, not just read about it.

Time

You have 72 hours from when you receive it to submit.

What we're looking at

- Does it work? End to end, not just pieces.
- Extraction quality - is the output accurate and well-structured, or did the agent just guess?
- RAG - does the knowledge base actually improve the result, or is it decoration?
- How you handle things going wrong - bad files, ambiguous data, missing fields.
- Architecture and code - clean, readable, sensible decisions.